

Student Housing Safety Assistant

Technical report and illustrated user manual

Maker: Bui Xuan Mai

Edition: 12 September 2026 · Logic review and learning examples

Git repository: https://github.com/S0lluxx26/Project_web_student_support

Live project: https://s0lluxx26.github.io/Project_web_student_support/

Demo/manual: https://s0lluxx26.github.io/Project_web_student_support/demo.html

Safety help and practice: https://s0lluxx26.github.io/Project_web_student_support/help.html

AI prompt reference: [docs/AI_PROMPTS.md](#)

This report explains what the project does, how its repository is organized, how data moves through the browser, and how to reproduce, test and deploy it. It combines an architectural review with the illustrated user manual, suggested AI prompts and evidence from the implemented system.

The central assessment uses curated rules. OCR reads Korean chat screenshots. The optional small LLM explains a selected finding on a PC. These components have different responsibilities and limitations.

All included chat fixtures are fictional. Application screenshots are actual captures. Measured demonstrations do not establish real-world fraud-detection accuracy. The report distinguishes implemented features from work that requires real data, domain review or an external account.

Contents

The PDF contents entries include page numbers and clickable links. The Markdown entries below jump to the corresponding chapters.

- 1. Project overview and scope 3
- 2. Repository architecture and file structure 4
- 3. Deployment and data flow 6
- 4. System sequence diagrams 8
- 5. AI prompt sources and working method 10
- 6. Local setup and interface implementation 11
- 7. Implementing explainable conversation checks 12
- 8. Screenshot OCR and the review workflow 13
- 9. Demonstration results 15
- 10. Document review and private export 17
- 11. Optional desktop AI and its limitations 19
- 12. Demo and report reproduction 21
- 13. Verification and deployment 23
- 14. Results, limitations and remaining work 24
- 15. References and glossary 25

Figures: repository-to-deployment flow; application data-flow graph; screenshot/review/export sequence; optional desktop AI sequence. Editable Mermaid sources are linked beside each figure.

1. Project overview and scope

This project supports students reviewing rental conversations in Korea. Its purpose is to identify known warning signals, show the relevant evidence and help the user decide what to verify next. It does not determine whether a named person or property is fraudulent.

Intended users and core tasks

- A student pastes a landlord or agent conversation and reviews the result.
- A student selects a Korean chat screenshot, corrects OCR text and confirms speakers before checking it.
- A student records document answers and reviews an export before copying, sharing or printing.
- A PC user may explicitly try an experimental local AI explanation of one existing warning.

Scope and acceptance

Requirement	Implemented behavior	Practical limit
Static deployment	One build for GitHub Pages and Vercel	Vercel account import remains optional
Mobile access	Text checker, OCR and user manual	Optional LLM is disabled on phones/tablets
Private input processing	Messages and images stay in the browser	Hosts still receive ordinary asset requests
Explainable findings	27 curated patterns with evidence	Heuristics, not calibrated fraud probabilities
AI assistance	Explicit PC opt-in, CPU inference	About 397 MB download; meaning errors observed
Reproducible demonstration	Two illustrated examples plus 12 new practice screenshots	Not an evaluation of real-world accuracy

Technology distinction. The rule analyzer produces the assessment. Tesseract performs screenshot OCR. Wllama runs the optional Qwen model. The project has no trained rental-scam classifier, application database, scraper or server-side inference endpoint.

Contract photos are a separate local preview; they do not enter screenshot OCR or the helper. Language/theme preferences and model caches may persist, but chat content is not intentionally stored. A model cache belongs to a site origin, so different hosts do not share it. No service worker is installed; local processing is not a promise of offline reload.

Reading route. Readers assessing the design can start with Chapters 2-5. Developers can follow Chapters 6-8 and 11-13. Users can follow Chapters 8-10 and the public Demo. Evidence and remaining work are in Chapter 14.

2. Repository architecture and file structure

The repository contains a vanilla HTML/CSS/JavaScript application, curated JSON data, vendored browser runtimes, a documentation set and test/build scripts. There is no server application to start in production.

Selected repository file tree

This tree emphasizes the files needed to understand and reproduce the project. It is not an exhaustive listing of every test or historical review note. Generated folders are labeled explicitly.

```
Project_web_student_support/
|-- index.html           # main application
|-- demo.html           # illustrated KO/EN user manual
|-- help.html           # self-check guidance and 12 examples
|-- INSTALL.md          # PC setup and AI installation prompt
|-- README.md / PLAN.md # overview and current work
|-- package.json / package-lock.json
|-- vercel.json / .vercelignore # optional static host configuration
|-- .github/workflows/deploy.yml # validation and Pages publication
|-- assets/
|   |-- css/            # style.css, manual.css
|   |-- js/
|     |-- app.js / i18n.js # boot, routes and language
|     |-- housing.js      # conversation UI and session state
|     |-- conversation.js  # message and speaker parsing
|     |-- analyzer.js     # rule assessment and evidence
|     |-- detector.js     # concept matching; similarity API
|     |-- ocr.js / redact.js # screenshot OCR and export masking
|     |-- llm.js / llm-ui.js / llm-worker.js
|     |-- goshiwon.js / manual.js
|     |-- theme.js / help.js # shared theme and practice page
|   |-- vendor/         # pinned OCR, Wllama and hash runtime
|   |-- examples/       # 12 PNGs, targets and OCR observations
|   |-- manual/         # fictional samples and actual captures
|     |-- cases.json    # recorded OCR/analysis outputs
|     |-- diagrams/     # generated SVG, PNG and manifest
|-- data/               # patterns, lexicon, documents,
|                       # examples, goshiwon and i18n JSON
|-- docs/
|   |-- AI_PROMPTS.md   # research and rule prompt playbook
|   |-- OCR.md / BROWSER_LLM_OPTIONS.md
|   |-- DEPLOY_VERCEL.md / IMPLEMENTATION_STATUS.md
|   |-- REVIEW_2026-09-12.md # logic and learning-example review
|   |-- FIELD_RESEARCH.md / REVIEW_ANALYSIS.md
|   |-- SCAM_PATTERNS.md / CONTRIBUTING.md
|   |-- benchmarks/    # recorded model measurements
|   |-- diagrams/      # editable Mermaid .mmd sources
|   |-- report-tools/   # optional locked Mermaid dependency
|-- manual/
|   |-- PROJECT_GUIDE.md # this report source
|   |-- project-guide.pdf # published PDF copy
|-- demo/               # original 22 fictional chat fixtures
|-- scripts/            # build, serve, tests, OCR/model checks
|   |-- capture-manual.mjs / render-report-diagrams.mjs
|   |-- build-guide-pdf.py / check-report.mjs
|-- dist/               # generated public website; Git-ignored
|-- output/pdf/         # generated PDF artifact; Git-ignored
|-- tmp/                # local scratch/results; Git-ignored
```

Module responsibilities

Component	Main files	Responsibility
Interface and navigation	index.html, assets/css/style.css, assets/js/app.js	Home, guided steps, language, loading and recovery
Data	data/patterns.json, lexicon.json, documents.json, i18n.json	27 patterns, concept rules, questions and KO/EN text
Conversation checks	conversation.js, analyzer.js, housing.js	Speakers, demands, questions, evidence and assessment
Screenshot reader	ocr.js, assets/vendor/tesseract/	Local Korean recognition, progress and editable review
Export	redact.js, housing.js	Editable masking preview before copy, share or print
Optional explanation	llm.js, llm-ui.js, llm-worker.js	Desktop opt-in, verified model, worker and draft checks
Demo and guide	demo.html, assets/manual/, manual/	Fictional samples, actual captures and downloadable guide
Safety help/theme	help.html, help.js, theme.js, assets/examples/	Independent checks, 12 examples and a shared viewing preference
Delivery	scripts/build-vercel.mjs, deploy.yml, vercel.json	One artifact for GitHub Pages and optional Vercel

JavaScript filenames in the table are under assets/js/ unless another path is shown. The deployment workflow is `.github/workflows/deploy.yml`.

Important integration detail. `analyzer.js` calls `Conversation.parse()` and `Detector.matchRules()`. The near-duplicate `Detector.scan()` API has no application consumer; it should not be presented as a shipped review-analysis feature.

Repository versus publication. `scripts/build-vercel.mjs` stages an allowlist into `dist/`: three HTML pages, the published guide, assets, JSON and license files. The complete `docs/` and `scripts/` trees, original test fixtures, `node_modules`, model GGUF files and private keys are not website assets. Selected fictional samples, captures and recorded observations under `assets/manual/` and `assets/examples/` are public. Report-only Mermaid dependencies stay out of the ordinary application installation.

3. Deployment and data flow

The following diagrams are generated from committed Mermaid sources. Solid arrows show the principal processing path. Labels on dotted arrows describe supporting inputs or optional actions. Human review is an explicit step.

3.1 Repository to deployment




Figure 1. Actions validates before GitHub Pages publication. Vercel is an optional separate account connection. Its configuration is supplied; no Vercel site is claimed live.

[Editable Mermaid source](#) · [Full-size diagram](#)

Both hosts serve static files. The visitor runs OCR and the optional LLM. Moving between hosts does not transfer inference work to a cloud server. The developer laptop does not need to remain online after publication.

3.2 Application data-flow graph




Figure 2. Local processing and external boundaries. The browser receives assets and model weights; the application does not send the user conversation or images to those sources.

Editable Mermaid source · Full-size diagram

Flow	Payload and destination	Control
Screenshot to review	Image bytes to local OCR; draft back to UI	User corrects and applies before analysis
Review to analyzer	Reviewed text, speaker/context input and manual answers	Unknown answers remain unknown
Analyzer to report	Matched rules, evidence, assessment and questions	A match is a warning signal, not proof of fraud
Report to helper	One selected rule guidance and fixed assessment	PC opt-in plus explicit model load
Model source to worker	Exact local file, verified cache or pinned download	Size/SHA-256 validation; no transcript upload
Report to export	Masked, editable report text	Final user confirmation is required

4. System sequence diagrams

These sequence diagrams describe the current code paths rather than a proposed backend. A worker is a browser background execution context, not a server. The cancel branches matter because an old result must not reappear after a reset.

4.1 Screenshot, review, analysis and export




Figure 3. Recognition stops at an editable draft. The user must apply it before rules run, and must review another preview before export.

Editable Mermaid source · Full-size diagram

The UI checks file limits. Session identity prevents stale OCR callbacks after reset. Bubble-side selection identifies your messages; confirm the other speaker's role separately.

4.2 Optional desktop AI explanation




Figure 4. Model loading and explanation are explicit actions. Model source means the selected local file, the cache for this origin or the pinned Hugging Face download.

Editable Mermaid source · Full-size diagram

The model source receives a request for weights only when downloading is necessary. A local-file choice avoids that download. The optional worker verifies the bytes and uses CPU inference. Validation can reject malformed output but cannot guarantee the meaning of an accepted draft. Cancel/reset terminates the worker and clears model state.

5. AI prompt sources and working method

Primary prompt reference:

https://github.com/S0lluxx26/Project_web_student_support/blob/main/docs/AI_PROMPTS.md

The docs/ playbook is the existing source for research and rule-authoring prompts. It covers investigating tactics, turning an investigated tactic into a rule, generating paraphrases and benign controls, finding Korean terminology and challenging a rule before release. Its opening has been updated to acknowledge OCR and the optional runtime LLM.

The report does not claim that all prompts below were copied from that playbook. Its installation and implementation templates were written for the current system. Neither file is authenticated evidence of the exact prompts used during past development.

Purpose	Read first	Expected output and review
Investigate a warning tactic	docs/AI_PROMPTS.md, section 2.1	Candidate mechanisms and ordinary comparisons; verify sources
Draft or challenge a rule	docs/AI_PROMPTS.md, sections 2.2 and 5	Proposed concepts, benign controls and failure cases
Extend paraphrase coverage	docs/AI_PROMPTS.md, section 2.3	Synthetic regression examples, not a real evaluation set
Install and run on a PC	INSTALL.md, section 1	A running local site and actual verification results
Implement UI, OCR, export or AI	Suggested prompts in Chapters 6-8 and 10-12	Reviewable code and relevant tests
Understand runtime model input	BROWSER_LLM_OPTIONS.md and llm.js	System prompt, structured finding guidance and output checks

How to use these sources with an AI assistant

1. Give the assistant the repository URL and ask it to read README.md, PLAN.md and the relevant docs/ file before editing.
2. State one bounded task, identify the files involved and describe the desired user behavior.
3. Ask for both suspicious and ordinary examples when changing detection logic.
4. Ask the assistant to identify unsupported claims instead of inventing sources, evaluation data or benchmark results.
5. Review the diff and run the relevant checks. Treat generated text as a draft until its facts and behavior are verified.

Suggested starting prompt

Read this repository and its documentation before proposing changes:

https://github.com/S0lluxx26/Project_web_student_support

Start with README.md, PLAN.md and docs/AI_PROMPTS.md.

For PC setup also follow INSTALL.md. For runtime AI behavior read docs/BROWSER_LLM_OPTIONS.md and assets/js/llm.js.

Identify which features exist in code and which remain proposals.

Keep the rule assessment separate from the optional explanation.

Make one reviewable change, test it and report actual outcomes.

Maker attribution: Bui Xuan Mai. Keep this credit and the repository address in derivative project documentation. Do not invent personal history or research results.

6. Local setup and interface implementation

Install Git and Node.js 22 or later. For a complete PC setup prompt and verification checklist, see [INSTALL.md](#). From a terminal:

```
git clone https://github.com/S01luxx26/Project_web_student_support.git
cd Project_web_student_support
npm ci --ignore-scripts
npx playwright install chromium
npm run build
npm run serve
```

Open <http://127.0.0.1:8765/> in the browser. Double-clicking `index.html` does not provide the HTTP environment needed by `fetch` and `workers`. Run `npm run verify:install` for the combined installation checks. Keep the server terminal running while using the local site; `Ctrl+C` stops it.

For development before rebuilding, use:

```
npm run serve -- --src
```


Suggested AI prompt

```
Inspect the repository before editing. Preserve its vanilla
HTML/CSS/JavaScript architecture and Korean/English support.
Make pasted rental messages and chat screenshots the two clear
entry points. Keep documents as supporting checks.
Add a visible Demo / guide link. Use accessible labels,
keyboard focus, readable contrast and mobile layouts.
Do not add a backend or automatically load an LLM.
After changes, report the files changed and actual checks run.
```

Implementation steps. Keep the home page focused on the conversation. Synchronize the home and detailed text fields. On screenshot entry, reveal the conversation screen before opening its OCR panel. Invalidate an old report when input changes. Allow the navigation actions to wrap on narrow screens.

Implemented result. The landing page offers text and screenshot input, a Demo / guide link and supporting document/goshiwon tools. Screenshot analysis goes directly to the report after the user reviews and applies the draft.

The report has six numbered sections: summary, input coverage, evidence, next checks, reviewed export and optional AI. Evidence highlights and action accents help users locate the reason and next step. The header links to Safety help. A Dark mode toggle persists a viewing preference across the checker, Demo and help pages; it stores no chat content.



Actual home screen with text, screenshot and Demo entry points.

7. Implementing explainable conversation checks

The curated catalogue has 27 patterns. The analyzer considers text, speaker role, context and manually answered document questions. It distinguishes questions or denials from demands where supported. It displays evidence excerpts and suggested next actions.

The assessment states are strong warning signals, needs review, no known signals and insufficient information. These are project categories, not calibrated probabilities. A result with no known signal does not prove a person, listing or contract is safe.

Suggested AI prompt

Review conversation.js, analyzer.js, patterns.json and lexicon.json.
 For each proposed change, give a suspicious example and an ordinary example with similar words. Preserve unknown speakers and distinguish questions, refusals and reported speech.
 Return matched evidence and explain the next check.
 Do not output a percentage chance of fraud.
 Do not turn a missing document answer into a verified answer.
 Validate pattern IDs, category/severity values, concept references and bilingual keys. Identify unsupported legal claims separately.

Implementation steps. Parse messages with speaker labels; keep unrecognized roles unknown. Match curated phrases and concepts. Apply context and suppression rules. Combine manual answers without treating unknown as no. Render the assessment, evidence and questions. Add regression cases for ordinary wording as well as suspicious demands.

Implemented result. The release uses deterministic analysis and documented heuristics. Source-related references exist for all 27 patterns, but exact legal-claim review remains incomplete. The similarity engine in detector.js is available as a library; a review-scraping interface has not been implemented.

Logic review in this edition

The analyzer now preserves sentence ownership directly and checks each candidate before deduplicating a pattern. A harmless first occurrence cannot hide a later demand, and identical words from two speakers keep their separate ownership. Speech acts apply to the quoted sentence, so a question at the end of a message does not erase an earlier demand.

Explicit polite payment requests are distinguished from ordinary document questions. Short role aliases no longer match inside personal names. Consent questions and explicit delayed-consent clauses cannot suppress a trust-property warning. A narrow pre-viewing payment rule has paired no-payment and risky-conditional controls.

The [review record](#) maps each defect to a regression test. These fixes improve the tested behavior; they do not establish accuracy on real conversations.

Data collection decision

This workflow does not require collecting Jikbang listings or website reviews. Rental messages supplied by the user already match the central task. A listing database would introduce a separate access, maintenance and evaluation problem.

For future research, establish permitted access and usefulness first. Do not assume a site grants scraping permission. For a trained classifier, obtain consented conversations, remove identifiers, define annotation rules, use independent review and hold out a test set. Synthetic examples are useful regression fixtures but are not a substitute for a real evaluation corpus.

Result of this step: an explainable checker that can operate without training data, with its limits kept visible. Marketplace integration and a trained classifier remain separate future work.

8. Screenshot OCR and the review workflow

Use the Demo page to download fictional sample A. Its original repository file is `demo/screenshots/01-kakao-pressure.jpg`. The public sample is `assets/manual/sample-pressure.jpg`.

1. Open the checker in a new tab and choose the screenshot-reading entry point.
2. Select sample A in the file picker. Wait for the recognition progress to finish; Cancel stops a running operation.
3. Compare the original image, uncertain lines and the editable recognized text.
4. In this particular sample, select the right-hand side as your messages. For other images, inspect the layout first. The initial choice is "Don't label them".
5. Correct missing words, amounts or negations. Choose "Add and check now" only after reviewing.

Student Housing Safety Assistant
Check before you sign. A lost deposit is very hard to get back.

Safety help Dark mode Demo / guide Start over 한국어

The bubble positions can tell us who said what

☒ The right-hand side is me ☐ The left-hand side is me ☐ Don't label them

Recognised text (you can edit it)

Me: 2026년 9월 8일
Them: 안녕하세요어제문의주신원룸 건물로 연락드립니다
Me: 네 안녕하세요 이번 주말에 방 볼 수 있을까요
Them: 지금세입자가살고 있어서방을 못 보여드려요
Them: 사진으로충분히 보셨죠 상태 아주 좋습니다
Them: 오늘 가계약금 100만원만 먼저 보내주시면 방 잡아드릴게요
Them: 다른분도지금보고 계셔서 오늘 안에 결정하셔야 해요
Me: 등기부등본 먼저 보여주실 수 있나요
Them: 그건 계약하실 때 보여드립니다 걱정 안 하셔도 돼요
Them: 제명의 계좌로 보내주세요 사장님이 지방에 계셔서요

Show what was removed

Add and check now Just add to the conversation Discard

Conversation

Actual editable OCR draft after selecting the right-hand side in sample A.

Side selection separates "Me" from "Them"; it does not prove whether the other person is an owner, agent or manager. Confirm that role separately in the report. The captured OCR draft still includes a date line and imperfect spacing. No manual text corrections were inserted into the recorded run.

Suggested AI prompt

Use the vendored Tesseract.js worker for Korean chat screenshots.
Show progress, cancel and an editable review before analysis.
Keep speaker-side selection unknown until the user chooses.
Preserve edits across language changes and navigation.
Test cancellation during startup and image encoding, not only after recognition starts. Terminate workers and discard stale results on reset. Use real OCR fixtures as well as mock tests.

Implemented result. Tesseract.js 5.1.1 runs locally using the vendored Korean model. The reader combines two segmentation attempts and presents a draft. It supports up to 5 images, 8 MiB per file, 25 MiB combined and 12 million decoded pixels per image. It does not upload screenshots.


Limits. Only Korean OCR is included. Mixed-language text, distorted captures and small fonts may fail. Real phone screenshots with consent still need evaluation. Contract photos do not use this OCR flow.

9. Demonstration results

The public capture record is `assets/manual/cases.json`. It contains fixture hashes, the recognized draft, selected side, browser version and actual matched signals. The capture script used real OCR and rule analysis without a mocked engine.

Sample A: pressure

Observed result: `strong_warning_signals`, 4 scored signals. The chat combines pressure to send money, an agent's account request, refusal to show the room and delayed document access. Inspect the actual quoted evidence before deciding what to ask next.




Actual result for the fictional pressure conversation: four warning signals need attention.

Sample B: ordinary

Reset the session before selecting sample B so unrelated conversations do not accumulate. The original file is `demo/screenshots/02-kakao-ordinary.jpg`.

Observed result: `no_known_signals`, 0 scored signals. This means the current rules did not find a known signal in the supplied text. It is not a verified safe contract.



Actual result for the ordinary fictional conversation: no known signals.

Result of this step: a side-by-side teaching example of differing system behavior. These two examples do not measure precision, recall or fraud detection accuracy.

Twelve additional practice screenshots

Open [Safety help](#) and [practice](#) for independent checks, category filters, downloadable PNGs and matching transcripts. Reset between unrelated examples. Compare speaker labels, amounts and negations before applying OCR text.

Teaching category	Three fictional cases	Expected reviewed result
Ordinary	Documents available; tenant refusal; no payment pressure	No known signals
Careful/uncertain	Short input; substitute-room pressure; reported restriction	Insufficient information or needs review
Pressure	Polite payment request; later demand; payment before viewing	Strong warning signals
Strong risk indicators	Third-party payment; registration restriction; delayed trust consent	Strong warning signals

The expected outcomes were specified in `assets/examples/cases.json`. Actual drafts and outcomes are recorded separately in `assets/examples/observations.json`. All 12 reached their intended assessment and required findings without text corrections in the recorded real-browser OCR run. Some headings and footers contain recognition errors. Three screenshots use dark chat backgrounds.

These are synthetic regression and teaching examples. They are not real phone captures, proof of fraud or a real-world accuracy measurement. The help page explains why ordinary behavior and no known signals still require independent checks, with links to official HUG and government guidance.

10. Document review and private export

Open the document-check route to read the requested checks. Answer only what has actually been verified. The contract step allows local image preview beside its questions. There is no automated contract OCR, authenticity check or AI contract verdict.

Student Housing Safety Assistant
Check before you sign. A lost deposit is very hard to get back.

Safety help Dark mode Demo / guide Start over 한국어

Independent housing safety check

Documents Conversation **Contract** Result

Contract check (optional)

Upload a photo of the contract to check it clause by clause side by side. The photo stays in your browser and is never sent anywhere. There is no automatic OCR — reading it yourself is the most reliable check.

Upload a photo of the contract

Choose Files No file chosen

What to check on the contract — tick anything that applies


Does the landlord's name differ from the registered owner?
Compare the name and birth digits in the landlord field against section 갑구 of the register, character by character.

☐ Yes ☐ No ☐ N/A ☒ Don't know

Are the deposit and rent amounts not written clearly in both words and figures?
Amounts should appear in words and figures together, which makes later alteration much harder.

Actual contract checklist screen with no real contract or personal information.

Copy, share and print open an editable preview. Automatically detected phone numbers, account-like numbers, email addresses and other supported identifiers are masked. Pattern matching can miss personal names or uncommon formats. Read and edit the preview before pressing the final action button.



Actual editable export preview for fictional sample A.

Suggested AI prompt

Preserve unknown document answers. Keep contract images local and clearly label them as previews, with no automated verdict. Route copy, share and print through an editable redaction preview. Show which identifier categories were masked and retain a deliberate restore-original control. Never claim perfect masking. Invalidate stale export text when input changes. Verify that direct browser printing cannot expose raw housing evidence.

Implemented result. Manual checks remain separate from OCR. Copy/share/print use the reviewed preview. Direct browser printing of the housing page asks the user to use that preview. Model drafts are excluded from report exports. Reset clears inputs, image previews, results and model state.

11. Optional desktop AI and its limitations

Open a report containing warnings, expand the AI explanation panel and read its limitations. On an eligible laptop or PC, enable the experiment. That checkbox alone starts no model download.

Then choose Download / load cached model, or select the exact previously downloaded GGUF file. Once ready, select one warning and generate a draft. Cancel stops work; Unload releases model memory; Remove cached model deletes this application's saved weights on the current site origin.

Yes: the model runs from the downloaded or selected file bytes on this device. Later loads reuse a verified browser cache when available. Browser eviction or removal can require another download; a different site origin has separate storage. The panel identifies whether the active model came from the network, verified cache or local file.

Phones and tablets cannot enable the helper. Detection uses browser device hints, mobile/tablet user agents and the touch-capable iPad desktop-mode signature. A narrow PC window does not disable it. This is best-effort device detection, not a guarantee of hardware performance. Text checking and OCR remain available on mobile.

Student Housing Safety Assistant
Check before you sign. A lost deposit is very hard to get back.

Safety help Dark mode Demo / guide Start over 한국어

▼ 6 Try an on-device AI explanation (optional)

Performance experiment. This model produced inaccurate explanations in testing. Use the reviewed report above and official sources when making decisions. The model cannot change the assessment.

Optional on laptops / PCs. No GPU required. About 397 MB to download; around 1.7 GB browser memory and 24–62 seconds per explanation on the tested desktop. Your device may be slower or unable to run it.

1. Read the limits and opt in

☒ Enable this experiment on my device (no automatic download)

2. Load the model

The downloaded model file is saved in this browser and runs on your device. Later loads reuse the verified cache. If the browser removes it, a download is needed again. You can also select the exact GGUF file yourself.

Qwen3 0.6B · Q4_K_M · 396.7 MB

Model files download from Hugging Face, which receives normal download request information. Your conversation and images are not sent. Inference runs in your browser. A different site address has a separate model cache.

Download / load cached model Or choose the downloaded GGUF file No file chosen

Use the Qwen3-0.6B-Q4_K_M.gguf file specified for this experiment. Its size and checksum are checked locally; selecting a file does not upload it.

3. Explain one warning

Actual desktop opt-in controls, before any model has been loaded.

The interface separates opt-in, load and generate into three stages. Before Generate explanation draft, a visible warning gives the measured wait and says a slower laptop may take several minutes or fail. The busy status repeats that warning. Users can continue reading the rule report or cancel.

Student Housing Safety Assistant
Check before you sign. A lost deposit is very hard to get back.

Safety help Dark mode Demo / guide Start over 한국어

Allow time: explanations took 24–62 seconds on the tested desktop. A laptop may take several minutes or fail. You can read the original report while waiting and cancel at any time.

Choose one warning to explain

Pressure to pay immediately

Generate explanation draft

Load the model only if you want to try a draft.

Model memory and storage

Remove this site's cached model

This site is informational and is not legal advice. Consult a professional about your specific situation.

Where to get help
Police 112 - cybercrime ecrm.police.go.kr
Free legal advice: Korea Legal Aid Corporation 132
HUG lease-fraud support centre 1566-9009

Actual generation wait warning and controls before a model is loaded.

Runtime and measured cost

- Wllama 3.6.1 runs llama.cpp through WebAssembly in a worker.
- The pinned Qwen3 0.6B Q4_K_M file is 396,705,472 bytes, approximately 397 MB.
- File size and SHA-256 are checked before use, including local files and cached copies.
- This configuration sets GPU layers to zero. A GPU is not required.
- Default Pages and Vercel deployments use one CPU thread.
- A Windows i7-14700KF desktop with roughly 32 GiB RAM took 24.27-62.02 seconds across 20 synthetic explanations; median 34.89 seconds.
- A browser process-group snapshot after loading was approximately 1.72 GB working set. This is neither a peak measurement nor exact model-only RAM.
- Other laptops may be slower or fail. These are not measurements from a typical laptop or phone.

[Wllama's browser documentation](#) describes CPU inference and its optional GPU support. This project's selected configuration is CPU-only.

Model boundaries and implementation prompt

The helper explains one existing rule finding. The UI does not give it raw messages, screenshots or contract photos. Its output cannot replace the rule assessment. Schema checks reject malformed output, invented finding IDs and several other invalid forms, but do not prove factual correctness.

Suggested AI prompt

Keep the rule report authoritative. Expose an optional desktop AI panel, disabled on phones and tablets. Do not use viewport width as the only device signal and do not require a GPU. Show model size, measured memory/latency and quality limitations before opt-in. Download nothing until the explicit load action. Use the pinned model and verify its size and SHA-256. Pass one selected rule's guidance, not the user's transcript. Render drafts as text, keep them out of exports, and preserve cancel/unload/cache removal. Test mobile blocking, desktop without WebGPU, failures and reset. Do not invent benchmarks.

Implemented result. Desktop visitors can discover and opt into the experiment without a special URL. Mobile visitors receive a disabled control and an explanation. Starting OCR unloads the LLM; starting the LLM releases OCR resources. Failed generation leaves the rule report available.

Quality evidence

The existing 20-case run accepted 19 outputs structurally. That is not 95% accuracy. Human inspection found changed conditions, contradictory wording and copied placeholders among accepted answers. One Korean request returned English and was rejected.

Real inference was also exercised on the public Pages origin: one recorded run loaded a local GGUF in 3.45 seconds, generated in 44.86 seconds and canceled another generation. These measurements confirm execution, not correct advice.

The experiment therefore remains optional and visibly cautioned. A small model can generate fluent but incorrect explanations; selecting a larger model or moving to Vercel does not establish better results.

Evidence: [implementation](#) and [benchmark notes](#), including the committed records in [docs/benchmarks/](#).

Result of this step: a testable local explanation experiment with limits, rather than a claimed AI fraud verdict.

12. Demo and report reproduction

The header's Demo / guide link opens [demo.html](#). Its Korean/English walkthrough covers samples, opening the reader, review, results, documents, export and the desktop helper. It includes downloadable fictional images and links to this Markdown guide and its PDF.

The published captures contain only fictional data. Two existing sample images are copied into [assets/manual/](#); 12 additional fictional practice PNGs and their records live in [assets/examples/](#). Contract screenshots show the actual checklist without a real contract image. No internet contract photo is needed for this demonstration.

Suggested AI prompt

Create a bilingual Demo / guide page using the existing fictional KakaoTalk fixtures. Capture actual application screens through Playwright and the real OCR engine. Show both a warning case and an ordinary case, with clear steps and limitations. Do not fabricate screenshots or measured results. Do not use private chat records or unlicensed contract photos. Credit maker Bui Xuan Mai and include the repository address. Write a detailed Markdown guide with architecture, reproducible AI prompts, steps and observed results. Review it for factual accuracy and attribution before converting the same text to PDF.

Regenerate the materials

```
npm run assets:stamp
npm run demo:capture
```

Review every changed image and assets/manual/cases.json. Update this Markdown if the observed behavior or instructions changed. The English guide is the source of the PDF, and both include the full repository address.

For the 12 learning screenshots, use the separate generator and explicit recording command:

```
python scripts/make-learning-screenshots.py
npm run build
node scripts/smoke-learning.mjs --record
```

The generator needs Pillow and a Korean font; EXAMPLE_FONT overrides the Windows default. Review all images, actual OCR drafts and correction flags. Preserve independent expected targets when a test fails. Normal verification is npm run test:learning and does not rewrite observations.

Suggested learning-feature prompt

Read docs/AI_PROMPTS.md and the current review record first.
 Add fictional ordinary, uncertain, pressure and strong-warning examples, each with an independent expected result and explanation.
 Test screenshot OCR separately from reviewed-text rule logic.
 Record corrections honestly. Keep real records and identifiers out.
 Provide a bilingual help page teaching independent verification.
 Preserve no-known-signals versus verified-safe distinctions.
 Number the report sections and test both light and dark layouts.
 Explain local model cache reuse and long generation waits.
 Run the release checks and report limits without accuracy claims.

To regenerate the PDF using the committed diagram images, install Python with reportlab, pypdf and Pillow, then run:

```
python scripts/build-guide-pdf.py
```

The builder reads manual/PROJECT_GUIDE.md and its relative images. It writes output/pdf/project-guide.pdf and the identical public copy at manual/project-guide.pdf. On Windows it uses Malgun Gothic for Korean glyphs; other systems must supply suitable TrueType fonts through GUIDE_FONT and GUIDE_FONT_BOLD. It does not install fonts.

Review sequence: read Markdown, verify commands/links/attribution, build PDF, extract its text to check completeness, render every PDF page and inspect layout. Regenerate the site artifact after the guide changes.

Result of this step: a public user manual and a reproducible English project guide attributed to Bui Xuan Mai.

Regenerate Mermaid diagrams

The editable definitions are in docs/diagrams/. Mermaid is isolated in docs/report-tools/ so normal website installation does not install the renderer. Diagram generation requires Node.js 22.12 or later and the installed Playwright Chromium.

```
npm ci --prefix docs/report-tools --ignore-scripts
node scripts/render-report-diagrams.mjs
python scripts/build-guide-pdf.py
npm run build
npm run test:report
```

Update the diagram source and its matching Mermaid block in this Markdown together. The renderer writes SVG/PNG assets and a source/output hash manifest. The PDF builder embeds these renders; it does not fetch an online rendering service. Review every changed figure and every PDF page before publication. The PDF contents page and bookmarks are generated from the report headings.

13. Verification and deployment

Run the release checks before committing:

```
npm run assets:stamp
npm run build
npm run test:ocr
npm run test:browser
npm run test:report
```

The build runs the unit/data suites, stages only explicitly public files and verifies copied bytes. Real OCR tests use synthetic image fixtures. Browser tests exercise navigation, reset/cancellation, privacy, host URL layouts, optional LLM controls and the manual's resources. Mock LLM tests check integration behavior, not model quality.

For a model or runtime change, rerun a real benchmark separately and inspect meaning:

```
npm run bench:llm -- --model tmp/models/Qwen3-0.6B-Q4_K_M.gguf --cases 20
```

GitHub and GitHub Pages

1. Review git status and git diff. Stage only intended source, documentation and public assets.
2. Commit and push to main using an authorized Git account. Keep SSH private keys outside the repository and deployment artifacts.
3. In repository Settings > Pages, use GitHub Actions as the publishing source.
4. The workflow installs locked dependencies, builds, runs real OCR/browser checks and uploads dist/ only after validation succeeds.
5. Wait for the deployment job to succeed. Open the live home page, Demo, Markdown and PDF; verify the new controls and a screenshot-to-report flow.

```
git status --short
git diff --stat
git add <reviewed-files>
git commit -m "Describe the verified change"
git push origin main
```

The angle-bracket value is a placeholder for reviewed paths, not a literal command argument. The repository is https://github.com/S0lluxx26/Project_web_student_support.

Optional Vercel deployment

Import the same Git repository into Vercel. Keep the repository root as the root directory, use the Other/static framework setting and retain vercel.json. The checked-in configuration installs dependencies with npm ci, builds with node scripts/build-vercel.mjs and publishes dist/.

This provides a second static host. It does not create a server-side model or require a cloud AI key. Vercel deployment is available after an account/project is connected; no Vercel production URL is claimed here. [Vercel build configuration](#)

Result of this step: one validated artifact can be served from a domain root or the GitHub Pages repository sub-path. The default public project link is the Pages address on this guide's cover.

14. Results, limitations and remaining work

What can be demonstrated now

Step	Observable result	Evidence
Interface	Text, screenshot and Demo entry points	index.html and the home capture
Rules	27 curated patterns with evidence and assessment states	patterns.json and analyzer tests
OCR	Editable real Korean recognition before apply	capture record and real OCR tests
Pressure sample	4 scored signals	cases.json and pressure-result capture
Ordinary sample	0 scored signals	cases.json and ordinary-result capture
Learning set	12 additional screenshots reached intended outcomes through real OCR	assets/examples/observations.json; synthetic cases only
Reading/help	Six numbered sections, light/dark toggle, independent checks	help.html and browser layout checks
Documents/export	Manual checks and editable export preview	contract/export captures and browser checks
Optional AI	Explicit PC opt-in; mobile disabled; no GPU requirement	support policy, worker config and browser tests
Real model	Runs on Pages; meaning errors observed	committed model benchmark records
Delivery	Shared dist;/ GitHub active, Vercel configurable	GitHub workflow, build script and vercel.json
Documentation	Maker and repository in Markdown and PDF	this guide and its generated PDF

Work that is still open

- Evaluate OCR on consented real phone screenshots, with human-reviewed text and speaker labels.
- Measure performance and usability on representative laptops. Test phone OCR separately; the optional LLM is disabled there.
- Review the 27 patterns' legal wording and numerical heuristics against exact current official sources with an appropriate reviewer.
- Obtain consented labeled data before attempting a trained classifier. Split by conversation/source and report false positives as well as misses.
- Connect and verify a Vercel project only if a second host is wanted.
- Conduct field research before adding a marketplace integration. No Jikbang collection or field-study results are claimed.

Troubleshooting

The page fails when opened from disk: run the local HTTP server.

OCR is empty or wrong: use a clearer Korean screenshot or paste corrected text; inspect amounts and negations.

AI control is disabled: phone/tablet browsers are intentionally blocked; unsupported browser capabilities can also block it.

AI load fails: retry the connection or use the exact GGUF file; wrong size/hash files are rejected.

A second host downloads again: Pages and Vercel have different origin storage.

No warning appears: inspect what was actually supplied and verified; absence of a match is not proof of safety.

15. References and glossary

Primary technical references

- Project repository
- GitHub Pages: static hosting
- Vercel: configuring a build
- Tesseract.js documentation and source
- Wllama: browser llama.cpp bindings
- Pinned model repository
- Project OCR notes
- Project model measurements
- Current project plan

Implementation and capture review date: 12 September 2026. External documentation can change; pinned project files and recorded measurements describe this edition.

Prompt and installation documentation

- AI Prompt Playbook: docs/AI_PROMPTS.md
- PC installation and its AI setup prompt: INSTALL.md
- Mermaid data-flow syntax
- Mermaid sequence diagram syntax
- Editable project diagrams

Short glossary

Term	Meaning in this project
OCR	Extracting editable text from an image
Rule analyzer	Curated phrase/concept logic producing warning signals
LLM	A language model generating an experimental explanation
Worker	Browser background execution, separate from the interface thread
WASM	WebAssembly code running on the visitor device
OPFS	Origin-private browser file storage used for model cache
GGUF	The model weight file format selected for Wllama
CI	Automated checks run by GitHub Actions before publication
dist/	Generated public site artifact, not the full repository

Traceability. Baseline reviewed for this edition: commit 192b565. This update fixes rule context/ownership handling, adds 12 fictional screenshots and Safety help, organizes the report and AI controls, and adds theme support. The captured original samples still produce four and zero scored signals respectively. Diagrams describe browser processing and static hosts; no backend inference service was added.